

Advanced Programming in the UNIX Environment

Week 02, Segment 3:
`read(2), write(2), lseek(3)`

**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

read(2)

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t num);
```

Returns: number of bytes read; 0 on EOF, -1 on error

read begins reading at the current offset, and increments the offset by the number of bytes actually read.

There can be several cases where read returns fewer than the number of bytes requested. For example:

- EOF reached before requested number of bytes have been read
- reading from a network, buffering can cause delays in arrival of data
- record-oriented devices (magtape) may return data one record at a time
- interruption by a signal

write(2)

```
#include <unistd.h>
```

```
ssize_t write(int fd, void *buf, size_t num);
```

Returns: number of bytes written if OK; -1 on error

- write returns the number of bytes written
- For regular files, write begins writing at the current offset (unless O_APPEND has been specified, in which case the offset is first set to the end of the file).
- After the write, the offset is adjusted by the number of bytes actually written.

write(2)

Some manual pages note:

If the real user is not the super-user, then write() clears the set-user-id bit on a file. This prevents penetration of system security by a user who “captures” a writable set-user-id file owned by the super-user.

Think of specific examples for this behavior.

Write a program that attempts to exploit a scenario where write(2) does not clear the setuid bit, then verify that your evil plan will be foiled.

Trying to create './newfile' with O_RDONLY | O_CREAT...

'./newfile' created. File descriptor is: 4

```
-rw----- 1 jschauma users 0 Sep 2 01:27 newfile
```

Checking if './newfile' exists...

```
-rw----- 1 jschauma users 0 Sep 2 01:27 ./newfile
```

Trying to create './newfile' with O_RDONLY | O_CREAT | O_EXCL...

Unable to create './newfile': File exists

Closing failed: Bad file descriptor

Trying to open './openex.c' with O_RDONLY...

'./openex.c' opened. File descriptor is: 5

'./openex.c' closed again

Trying to open (non-existent) './nosuchfile' with O_RDONLY...

Unable to open './nosuchfile': No such file or directory

Copied 'openex.c' to 'newfile'.

```
-rw----- 1 jschauma users 3192 Sep 2 01:28 newfile
```

Trying to open './newfile' with O_RDONLY | **O_TRUNC**...

'./newfile' opened. File descriptor is: 5

'./newfile' truncated -- see 'ls -l newfile'

```
-rw----- 1 jschauma users 0 Sep 2 01:28 newfile
```

apue\$

lseek(2)

```
#include <sys/types.h>
```

```
#include <fcntl.h>
```

```
off_t lseek(int fd, off_t offset, int whence);
```

Returns: new offset if OK; -1 on error



lseek(2)

```
#include <sys/types.h>
```

```
#include <fcntl.h>
```

```
off_t lseek(int fd, off_t offset, int whence);
```

Returns: new offset if OK; -1 on error

The value of whence determines how offset is used:

- `SEEK_SET` bytes from the beginning of the file
- `SEEK_CUR` bytes from the current file position
- `SEEK_END` bytes from the end of the file

"Weird" things you can do using `lseek(2)`:

- seek to a negative offset
- seek 0 bytes from the current position
- seek past the end of the file

[\$ ssh apue

Last login: Thu Sep 3 22:53:40 2020 from 10.0.2.2

NetBSD 9.0 (GENERIC) #0: Fri Feb 14 00:06:28 UTC 2020

Welcome to NetBSD!

[apue\$ cd 02

[apue\$ vim lseek.c

[apue\$ cc lseek.c

[apue\$./a.out

seek OK

[apue\$./a.out <lseek.c

seek OK

[apue\$ cat lseek.c | ./a.out

cannot seek

[apue\$ mkfifo /tmp/fifo

[apue\$ ls -l /tmp/fifo

prw-r--r-- 1 jschauma wheel 0 Sep 3 22:55 /tmp/fifo

[apue\$./a.out </tmp/fifo

cannot seek

apue\$ █


```
[apue$ echo $(( 10240020 / 512 ))
```

```
20000
```

```
[apue$ df .
```

Filesystem	512-blocks	Used	Avail	%Cap	Mounted on
/dev/wd0a	30497436	7378524	21594044	25%	/

```
[apue$ echo $(( 7378524 - 7378428 ))
```

```
96
```

```
[apue$ ls -ls file.hole
```

```
96 -rw----- 1 jschauma users 10240020 Sep  4 03:10 file.hole
```

```
[apue$ hexdump -c file.hole
```

```
00000000  a  b  c  d  e  f  g  h  i  j  \0  \0  \0  \0  \0  \0
00000010  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0
```

```
*
```

```
09c4000  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  A  B  C  D  E  F
09c4010  G  H  I  J
```

```
09c4014
```

```
[apue$ cp file.hole file.nohole
```

```
[apue$ ls -l file*
```

```
-rw----- 1 jschauma users 10240020 Sep  4 03:10 file.hole
-rw----- 1 jschauma users 10240020 Sep  4 03:12 file.nohole
```

```
[apue$ ls -ls file.*
```

```
  96 -rw----- 1 jschauma users 10240020 Sep  4 03:10 file.hole
20064 -rw----- 1 jschauma users 10240020 Sep  4 03:12 file.nohole
```

```
[apue$ hexdump -c file.nohole
```

I/O Efficiency

In simple-cat.c from last week, we used:

```
15 #define BUFSIZE 32768
[...]  
26 while ((n = read(STDIN_FILENO, buf, BUFSIZE)) > 0) {  
27     if (write(STDOUT_FILENO, buf, n) != n) {  
28         fprintf(stderr, "Unable to write: %s\n", strerror(errno));  
30         exit(EXIT_FAILURE);  
31     }  
32 }
```

How/why did we pick 32768? What if we increased/decreased that number?

Let's create a benchmark test:

```
mkdir -p tmp
for n in $(jot 10); do
    echo "Creating file number $n...";
    dd if=/dev/urandom of=tmp/file$n bs=$(( 1024 * 1024)) count=100 2>/dev/null;
done

for n in 3145728 1048576 32768 16384 4096 1024 256 128 64 1; do
    cc -Wall -DBUFSIZE=$n simple-cat.c;
    i=$(( $i + 1 ));
    for j in $$$(jot 5); do
        /usr/bin/time -p ./a.out <tmp/file$i >tmp/out;
    done
done
```


BUFSIZE = 16384

0.47	real	0.01	user	0.34	sys
0.32	real	0.00	user	0.30	sys
0.31	real	0.00	user	0.26	sys
0.29	real	0.02	user	0.25	sys
0.30	real	0.00	user	0.27	sys

BUFSIZE = 4096

0.59	real	0.02	user	0.41	sys
0.43	real	0.01	user	0.36	sys
0.45	real	0.02	user	0.36	sys
0.45	real	0.00	user	0.41	sys
0.44	real	0.01	user	0.41	sys

BUFSIZE = 1024

0.87	real	0.02	user	0.72	sys
0.69	real	0.04	user	0.62	sys
0.68	real	0.00	user	0.62	sys
0.68	real	0.01	user	0.66	sys
0.70	real	0.05	user	0.63	sys

Conclusion

Compare the manual pages for `read(2)`, `write(2)`, and `lseek(2)` on different OS.

`lseek(STDIN_FILENO, 0, SEEK_CUR)` succeeds when connected to a terminal - what happens when you try to seek to the end or the beginning? Does that even make sense?

If you create a new file, write, say, 10 bytes of data, and then seek to the end of the file, where do you end up? Why?

Play around with the creation and handling of sparse files as well as repeat our benchmark test on different operating- and file systems.

In our next segment:

How can we visualize multiple processes accessing the same files simultaneously?

Are there additional considerations regarding atomicity when using `read(2)` and `write(2)`?

How does file descriptor redirection work in the shell?