

Advanced Programming in the UNIX Environment

Week 08, Segment 2: System V IPC

**Department of Computer Science
Stevens Institute of Technology**

Jan Schaumann

`jschauma@stevens.edu`

`https://stevens.netmeister.org/631/`

System V IPC

There are three types of *asynchronous* IPC originating from System V:

- Semaphores
- Shared Memory
- Message Queues

All three use IPC structures, referred to by an *identifier* and a *key*; all three are (necessarily) limited to communication between processes on one and the same host.

Since these structures are not known by name, special system calls (`msgget(2)`, `semop(2)`, `shmat(2)`, etc.) and special userland commands (`ipcrm(1)`, `ipcs(1)`, etc.) are necessary.

System V IPC: Semaphores

A semaphore is a counter used to provide access to a shared data object for multiple processes. To obtain a shared resource a process needs to do the following:

1. Test semaphore that controls the resource.
2. If value of semaphore > 0, decrement semaphore and use resource; increment semaphore when done.
3. If value == 0 sleep until value > 0

Semaphores are obtained using `semget(2)`, properties controlled using `semctl(2)`, operations on a semaphore performed using `semop(2)`.

System V IPC: Semaphores

```
Terminal — 130x24

Press return to unlock:
Unlocked.
jschauma@apue$ ipcs -s
IPC status from <running system> as of Mon Oct 19 01:38:29 2020
Semaphores:
T      ID      KEY      MODE      OWNER      GROUP
s    720896  704662389 --rw-rw-rw- jschauma  users

jschauma@apue$ ./a.out
Press return to lock:
Trying to lock...
Locked.
Press return to unlock:
Unlocked.
jschauma@apue$ ipcrm -s 720896
jschauma@apue$ ipcs -s
IPC status from <running system> as of Mon Oct 19 01:39:47 2020
Semaphores:
T      ID      KEY      MODE      OWNER      GROUP
s    720896  704662389 --rw-rw-rw- jschauma  users

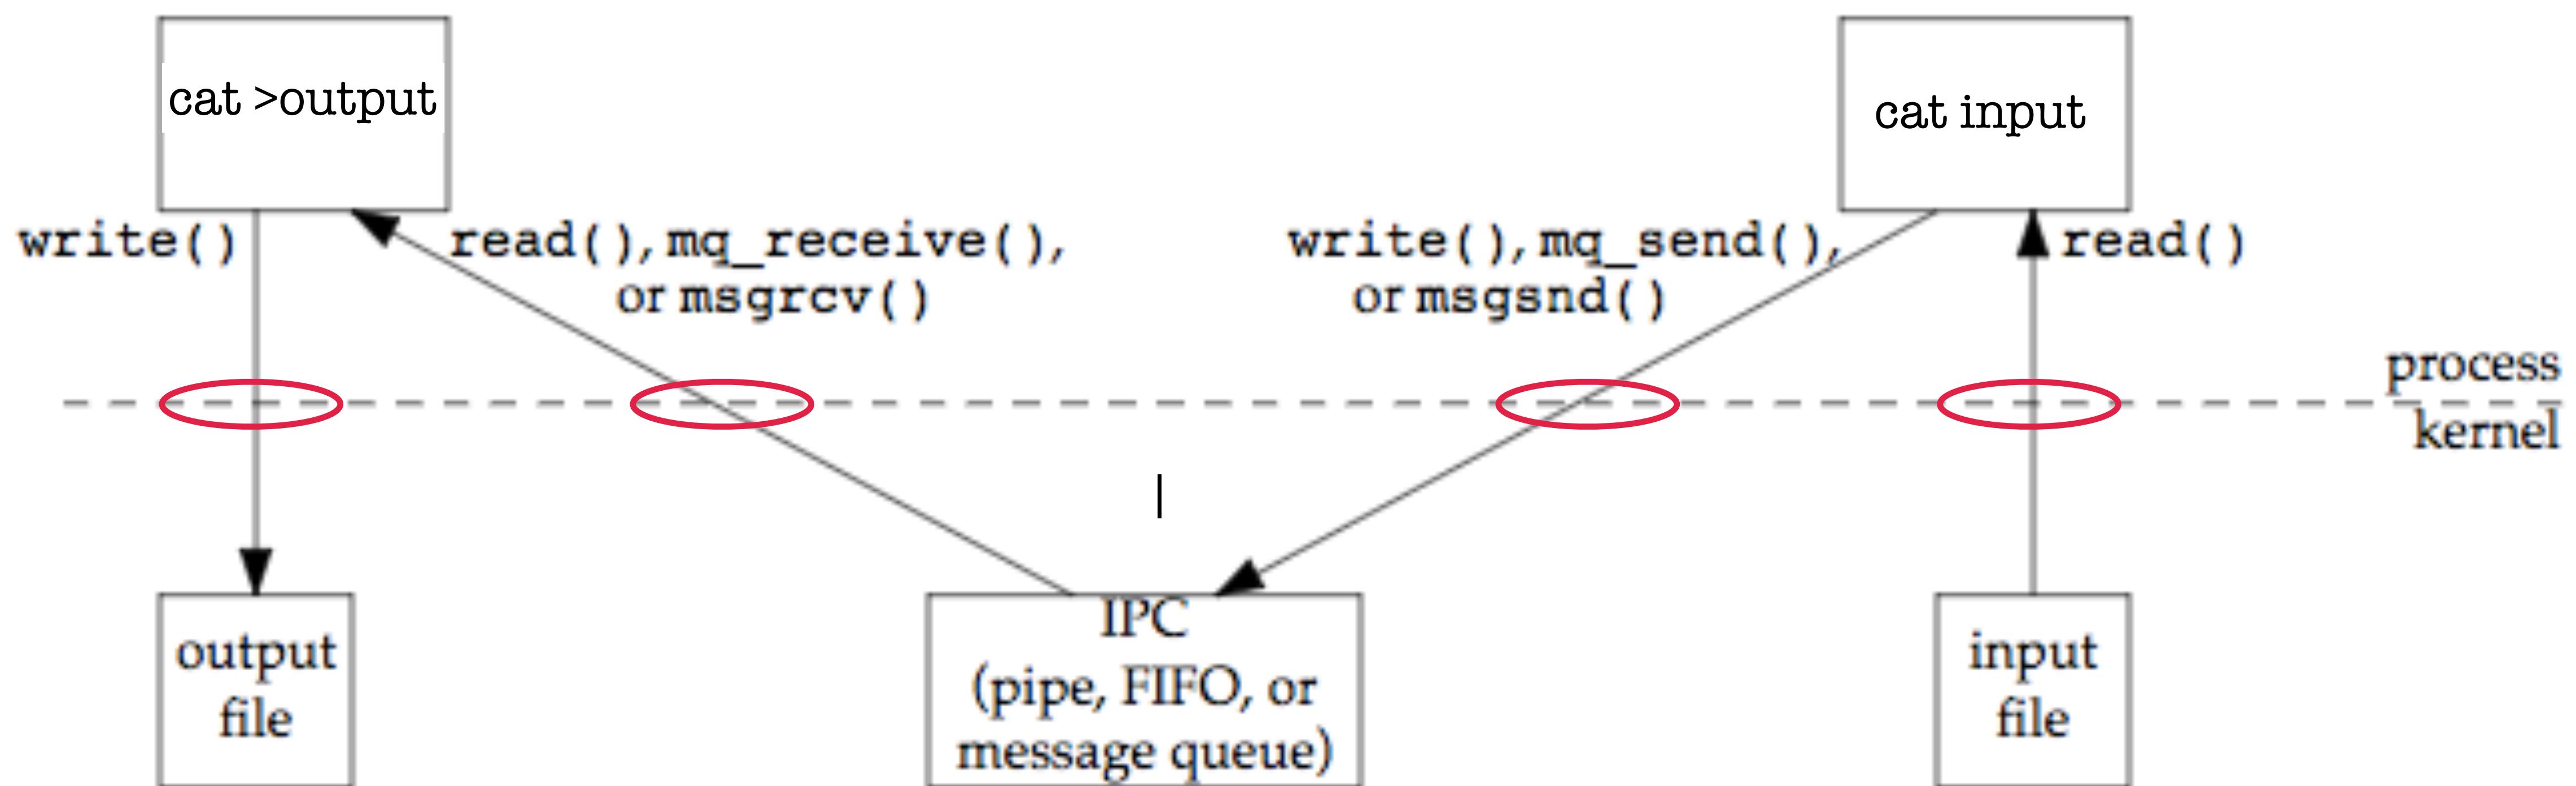
fred@apue$ ipcs -s
IPC status from <running system> as of Mon Oct 19 01:39:20 2020
Semaphores:
T      ID      KEY      MODE      OWNER      GROUP
s    720896  704662389 --rw-rw-rw- jschauma  users

fred@apue$ ./a.out
Press return to lock:
Trying to lock...
Locked.
Press return to unlock:
Unlocked.
fred@apue$ ipcs -s
IPC status from <running system> as of Mon Oct 19 01:39:20 2020
Semaphores:
T      ID      KEY      MODE      OWNER      GROUP
s    720896  704662389 --rw-rw-rw- jschauma  users

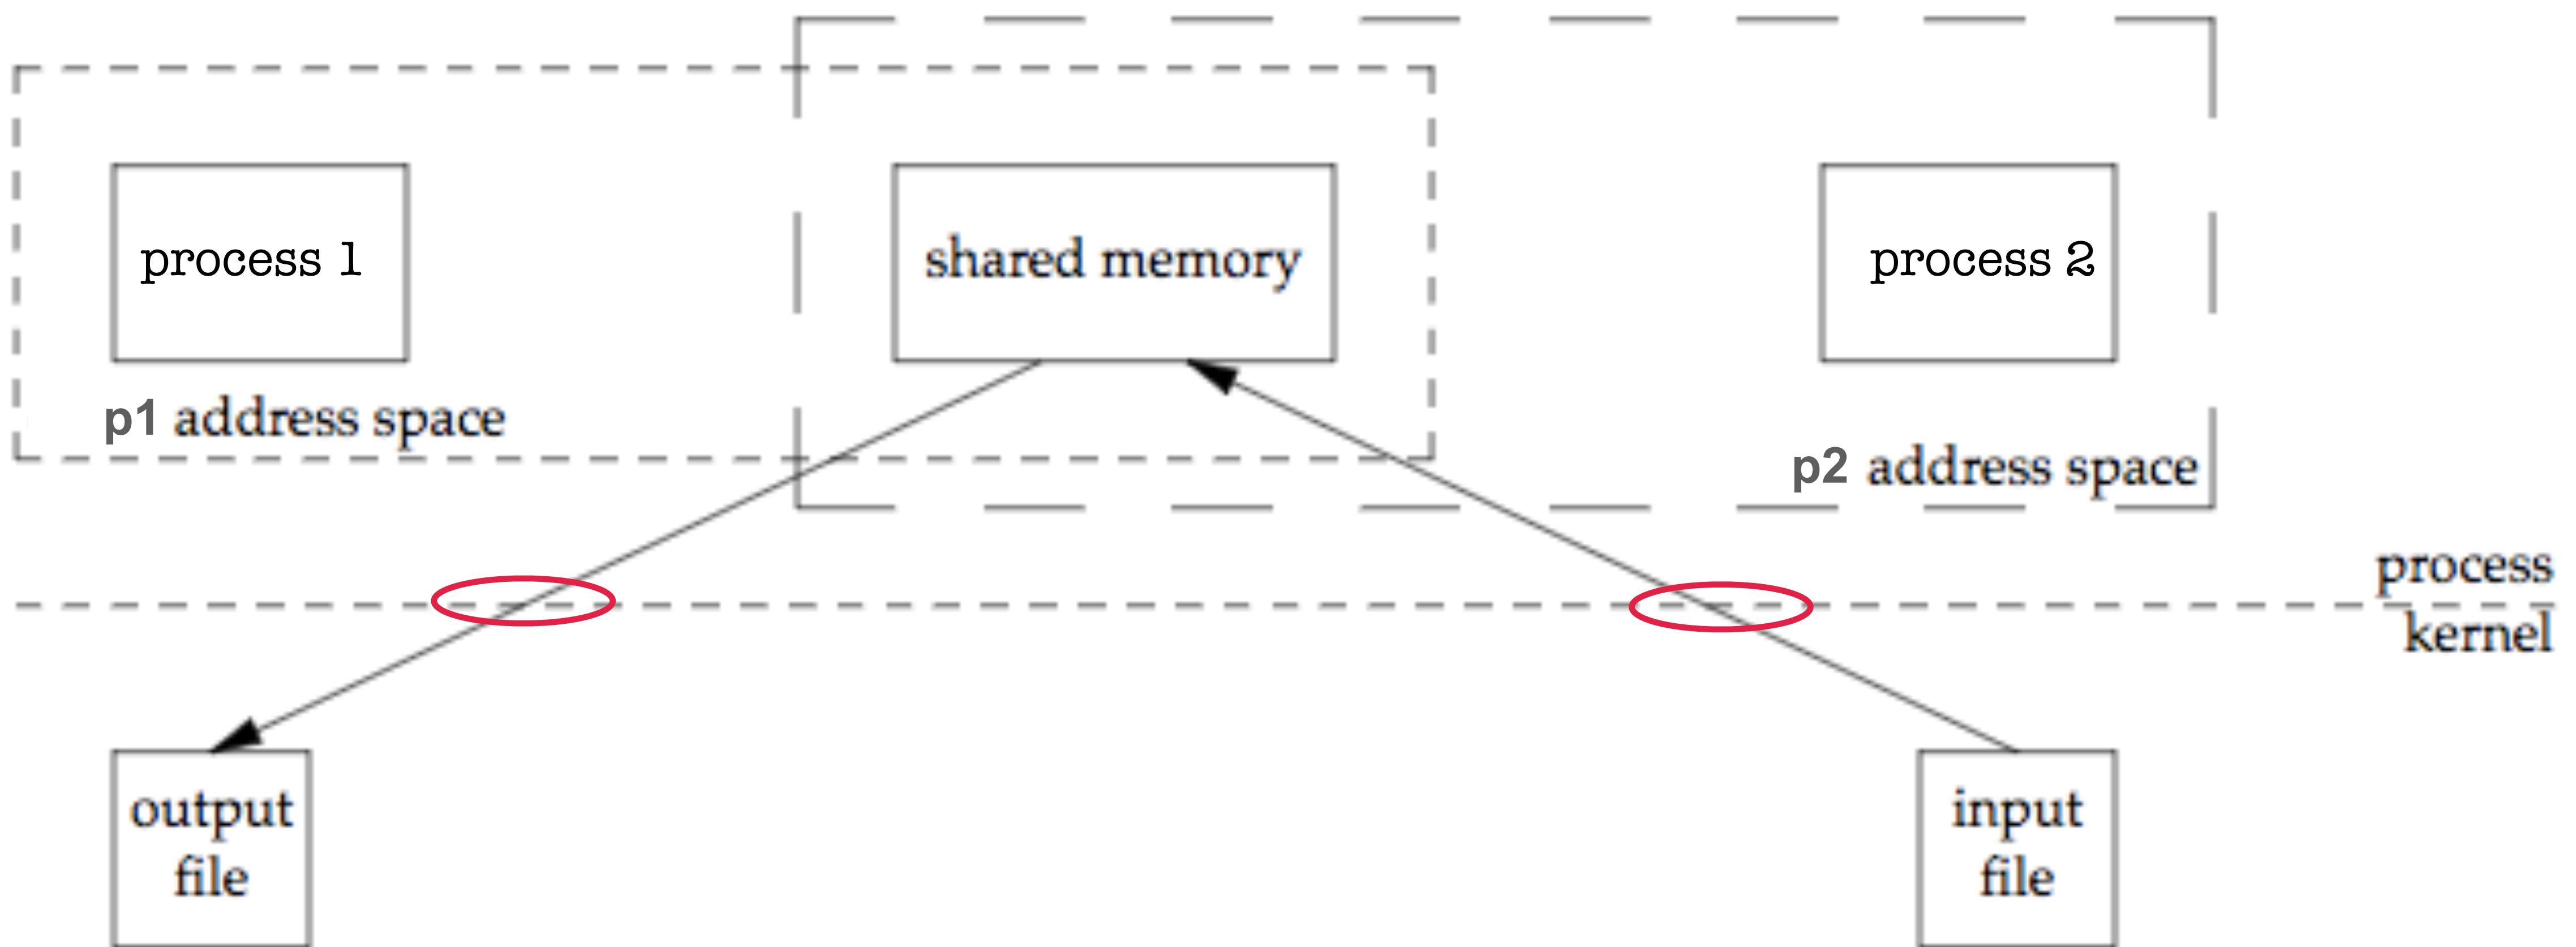
fred@apue$ ipcrm -s 720896
ipcrm: semid(720896): : Operation not permitted
fred@apue$ 1 sh
```

IPC Data Flow

```
$ cat input | cat >output
```



System V IPC: Shared Memory



System V IPC: Shared Memory

- fastest form of IPC
- access to shared region of memory often controlled using semaphores
- obtain a shared memory identifier using `shmget(2)`
- attach shared memory segment to a processes address space by calling `shmat(2)`
- detach it using `shmdt(2)`
- catch-all for shared memory operations: `shmctl(2)`

System V IPC: Shared Memory

```
m 327681 704662386 --rw-rw-rw- jschauma users jschauma users 0 1024 1074 1074 2:23:10 2:23:10 2:23:10
jschauma@apue$ ./a.out
1359: segment contains: "The cow says 'moo'."  

jschauma@apue$ ./a.out
1579: segment contains: "The cow says 'moo'."  

jschauma@apue$ ipcs -ma
IPC status from <running system> as of Mon Oct 19 02:23:52 2020

Shared Memory:
T      ID      KEY      MODE      OWNER      GROUP      CREATOR      CGROUP      NATTCH      SEGSZ      CPID      LPID      ATIME      DTIME      CTIME
m 327681 704662386 --rw-rw-rw- jschauma users jschauma users 0 1024 1074 1579 2:23:48 2:23:48 2:23:10

jschauma@apue$ ipcrm -m 327681
jschauma@apue$ ./a.out
1751: segment contains: ""
jschauma@apue$ ipcs -ma
IPC status from <running system> as of Mon Oct 19 02:25:06 2020

Shared Memory:
```

Terminal — 81x44

```
[jschauma@apue ./a.out
High address (args and env):
-----
envp[16] at : 0x7FFF2279D8
environ[16] at : 0x7FFF2279D8
envp[0] at : 0x7FFF227958
environ[0] at : 0x7FFF227958
last arg at : 0x7FFF227950
first arg at : 0x7FFF227948

Stack:
-----
First variable inside main at : 0x7FFF2278DC
func_array[] ends at : 0x7FFF2278D0
func_array[] (like 'array[]', but on stack) begins at : 0x7FFF2278C0
argc at : 0x7FFF2278BC
argv at : 0x7FFF2278B0
envp at : 0x7FFF2278A8
func2 (from main): frame at : 0x7FFF22788C
func frame at : 0x7FFF227884
static int n within func at : 0x 601EEC
func (called 1 times): frame at : 0x7FFF227884
func2 (from func): frame at : 0x7FFF22785C

Shared Memory:
-----
shared memory area ends at
shared memory area begins at : 0x7FF7EFE6A0
: 0x7FF7EE6000

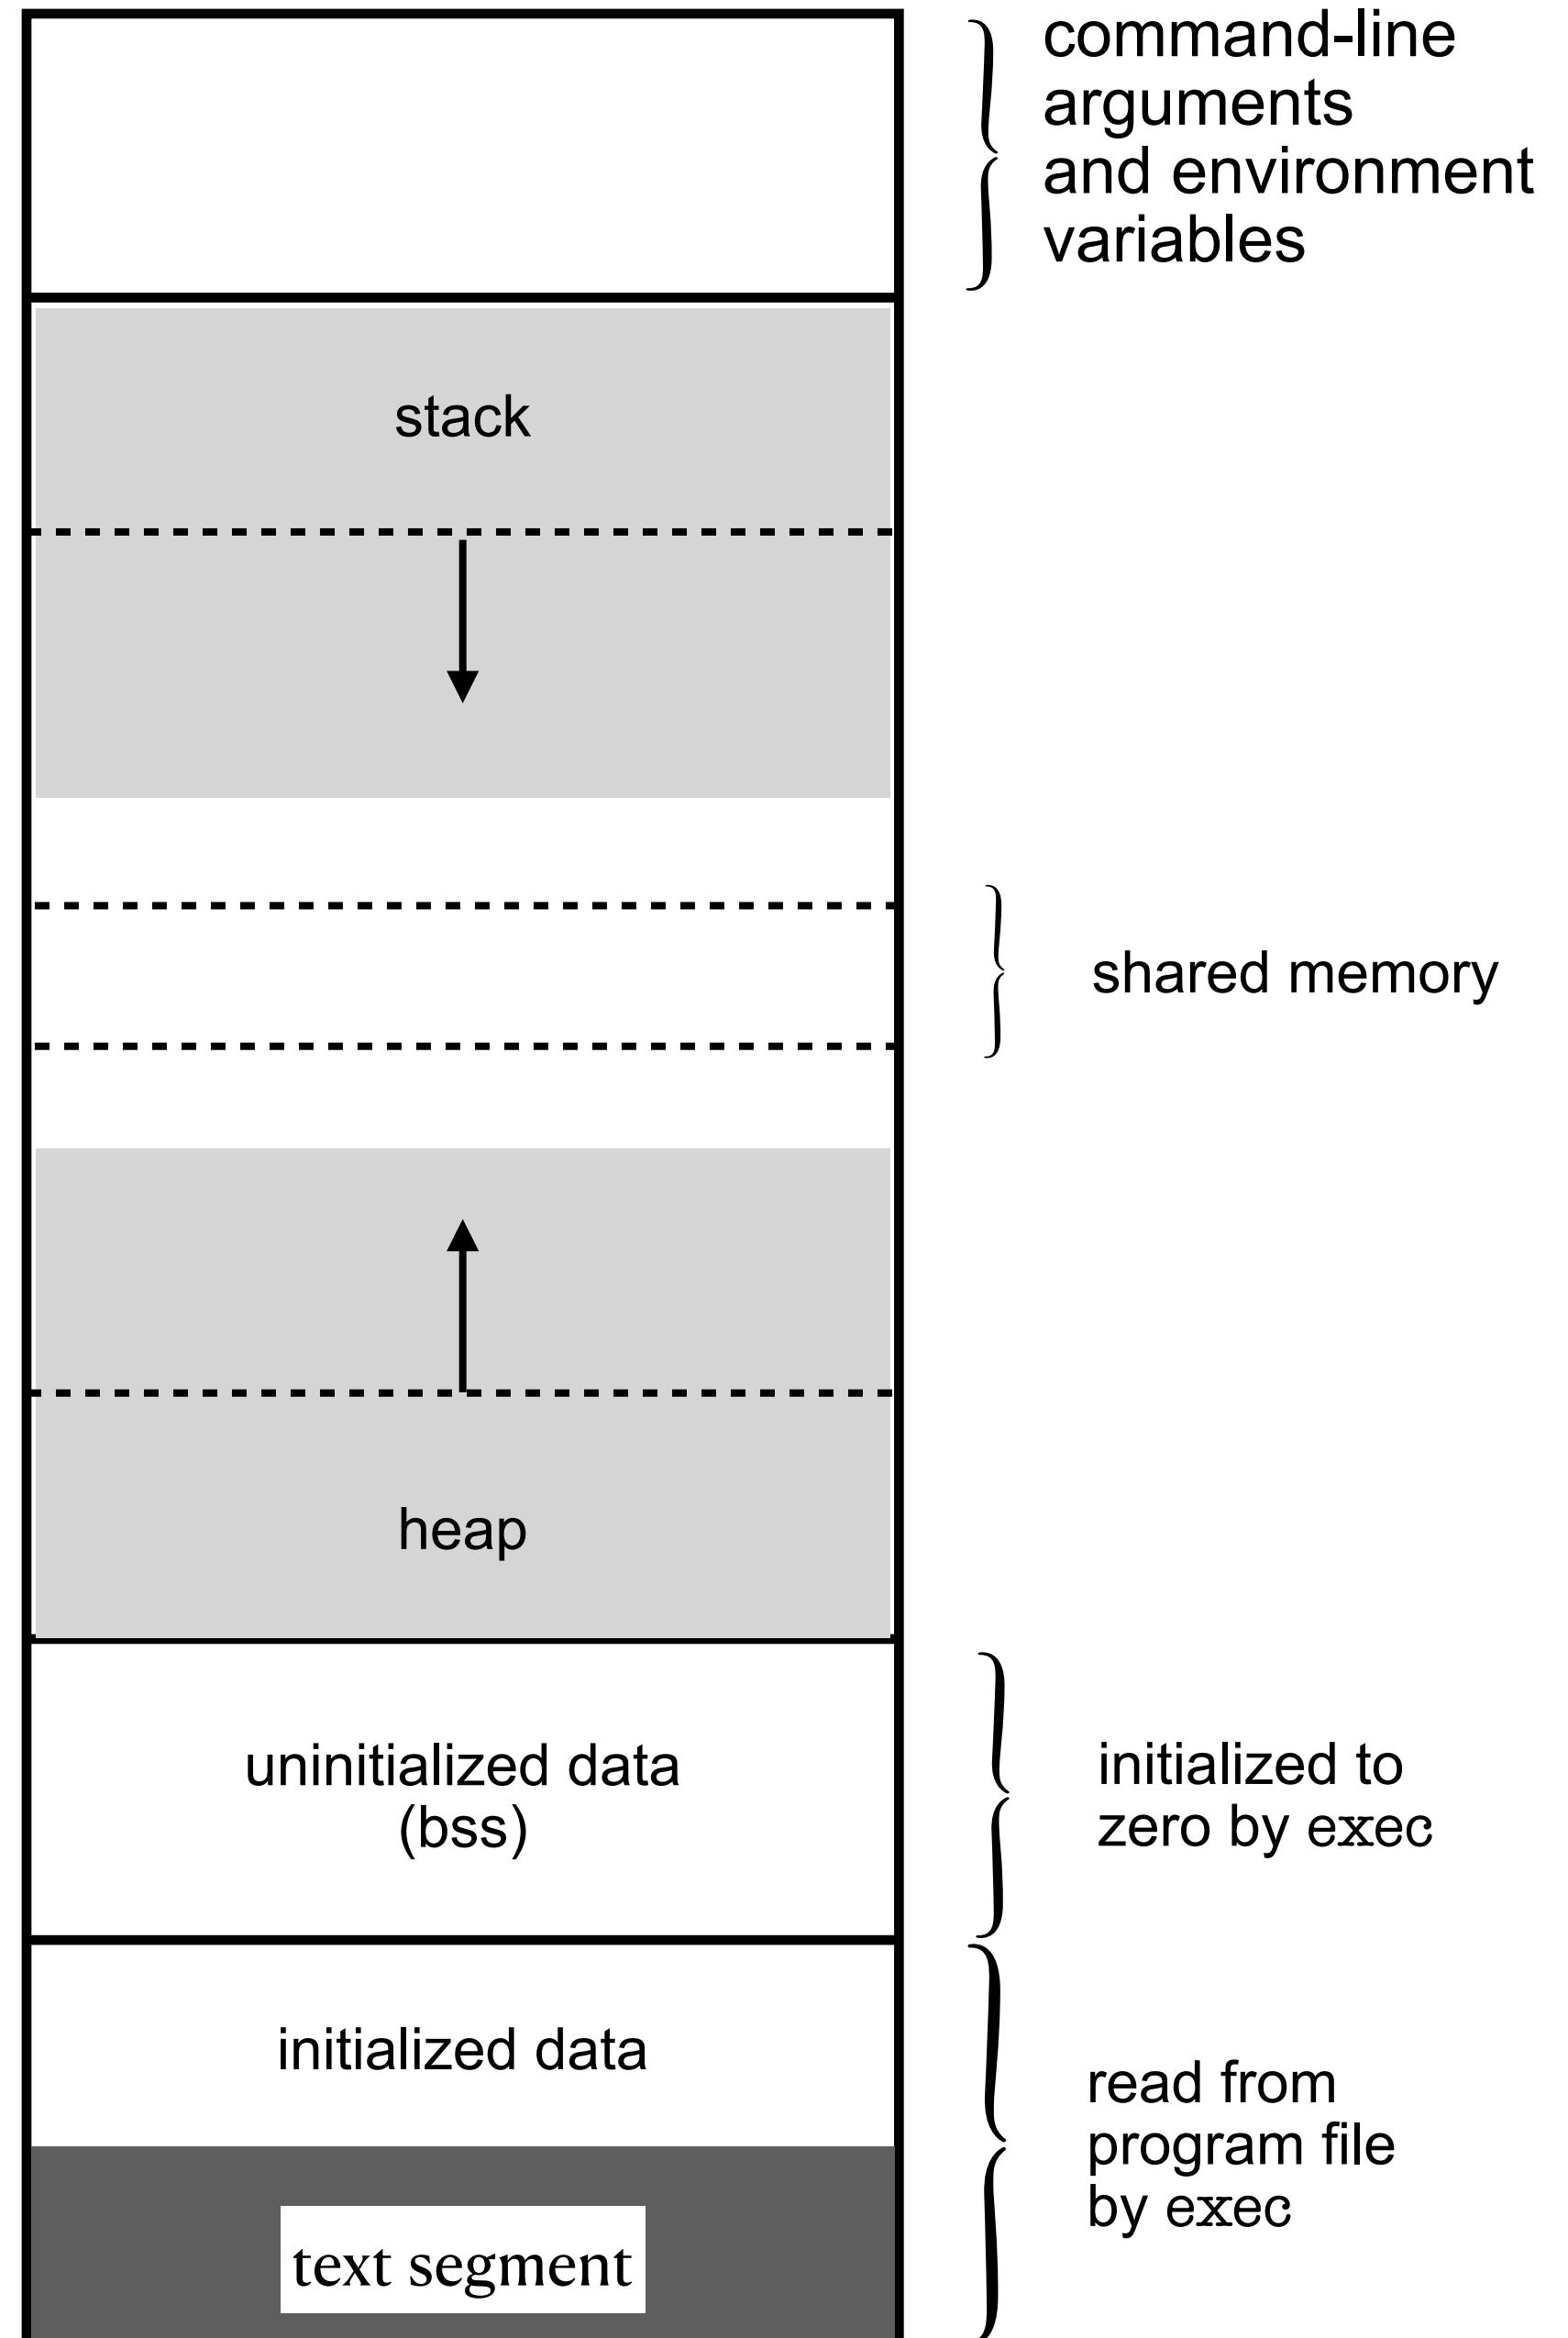
Heap:
-----
malloced area ends at : 0x7926986E5020
malloced area begins at : 0x7926986E5000

Uninitialized Data (BSS):
-----
array[] ends at : 0x 601F70
array[] (uninitialized, fixed-size char * on BSS) from : 0x 601F60
num2 (uninitialized global int) at : 0x 601F58
string2 (uninitialized global char *) at : 0x 601F50
extern **environ at : 0x 601F40

Initialized Data:
-----
```

high address

low address



System V IPC: Message Queues

- linked list of messages stored in kernel space
- create or open existing queue using `msgget(2)`
- add message at end of queue using `msgsnd(2)`
- receive messages from queue using `msgrcv(2)`
- control queue properties using `msgctl(2)`

The message itself is contained in a user-defined structure such as

```
struct mymsg {  
    long mtype;          /* message type */  
    char mtext[512];    /* body of message */  
};
```

System V IPC: Message Queues

```
Terminal — 130x24
q 393216      1 --rw-r--r-- jschauma    users jschauma    users      0      0 2048 1652 952 20:11:20 20:12:01 20:10:48

jschauma@apue$ ./rcv 1
Who will get this?
jschauma@apue$ ./send "Unblock thyself."
Usage: msgsend key message
jschauma@apue$ ./send 1 "Unblock thyself."
jschauma@apue$ ipcs -qa
IPC status from <running system> as of Mon Oct 19 20:13:28 2020

Message Queues:
T      ID      KEY      MODE      OWNER      GROUP      CREATOR      CGROUP      CBYTES      QNUM      QBYTES      LSPID      LRVID      STIME      RTIME      CTIME
q 393216      1 --rw-r--r-- jschauma    users jschauma    users      0      0 2048 225 1620 20:13:20 20:13:20 20:10:48

jschauma@apue$ ./rcv 2
msgget: No such file or directory
jschauma@apue$ ipcrm -q 393216
jschauma@apue$ ipcs -qa
IPC status from <running system> as of Mon Oct 19 20:13:48 2020

Message Queues:
T      ID      KEY      MODE      OWNER      GROUP      CREATOR      CGROUP      CBYTES      QNUM      QBYTES      LSPID      LRVID      STIME      RTIME      CTIME
jschauma@apue$
```

POSIX Message Queues

`mq(3)` provides a real-time IPC interface similar to System V message queues.
Notably:

- message queues are identified by a named identifier (no `ftok(3)` needed)
- message queues may or may not be exposed in the filesystem (e.g., `/dev/mqueue`)
- `mq_send(3)` and `mq_receive(3)` allow both blocking and non-blocking calls
- `mq_send(3)` lets you specify a priority; equal priority messages are queued as a FIFO, but higher priority messages are inserted before those of a lower priority
- `mq(3)` provides an asynchronous notification mechanism: `mq_notify(3)`

POSIX Message Queues

The terminal window shows three distinct sessions of message queue operations:

- Session 1 (Left):** A series of messages being enqueued. The output is:

```
Number of messages in queue: 1
Message of priority 1: TUNA
Message of priority 0: avocado
Message of priority 0: onion
Message of priority 0: tomato

Number of messages in queue: 1
Message of priority 0: avocado
Number of messages in queue: 1
Message of priority 0: onion
Number of messages in queue: 1
Message of priority 0: tomato
Number of messages in queue: 1
Message of priority 1: TUNA

Number of messages in queue: 1
Message of priority 0: avocado
Number of messages in queue: 1
Message of priority 0: onion
Number of messages in queue: 1
Message of priority 0: tomato
Number of messages in queue: 1
Message of priority 1: TUNA
```
- Session 2 (Middle):** A series of messages being sent to a queue. The output is:

```
0 sh
```
- Session 3 (Right):** A series of messages being received from a queue. The output is:

```
jschauma@apue$ ./mqsend avocado onion tomato
jschauma@apue$ ./mqsend -w avocado onion tomato
jschauma@apue$ ./mqsend avocado onion tomato
jschauma@apue$ ./mqsend avocado onion tomato
jschauma@apue$
```

System V IPC

- *asynchronous* IPC between processes on the same system
- old, but not obsolete
- semaphores are useful to guard access to shared resources
- shared memory allows for fast IPC
- message queues as a service are a popular way to implement "pub sub" models
 - Amazon Simple Queue Service
 - Apache Kafka
 - Java Message Service
 - RabbitMQ
 - ...